giftly

we help *agents* find products you *love*.

THE PROBLEM

agents are reading from *poisoned data*.

- × advertising
- × SEO & GEO
- × reviews
- × paid placement

30%

of online reviews are fake — and rising with AI bot spam.

275M

fake reviews Amazon let through in 2024 despite $500M and 8,000 employees.

83%

of consumers don't trust AI shopping agents — and they're right not to.

THE ANSWER

giftly is the trust layer.

the API AI agents query before they buy.

trygiftly.com

THE TRUST LAYER

the largest network of *human evaluators* on earth.

NETWORK
creators · experts · verified buyers

→

OUTPUT
honest, structured product data

→

CONSUMER
your shopping agent

every agent shopping with our data pays us.

today, our evaluators are *content creators*.

01 brands send creators free products

02 creators post only if they love them

03 creators earn affiliate commission on every sale

every match adds evaluation data to the long-term dataset.

in *1 week*, our network secured:

50
vetted creators

~2M
combined reach

7
official brand partnerships

$300M
combined partner valuation

THE BET

the future of agentic commerce will run on Giftly.

THE TRUST LAYER FOR AGENT COMMERCE

trygiftly.com

APPENDIX · 01 — WHY NOW

the infrastructure for buying is built. the infrastructure for *choosing* isn't.

AGENTIC COMMERCE SHIPPING IN 2026

- ChatGPT Instant Checkout
- OpenAI Operator
- Perplexity Shopping
- Claude + MCP
- Amazon Rufus
- Gemini Shopping

THE MARKET SIGNAL

50M ChatGPT shopping queries / day

1M+ Shopify merchants onboarding (Glossier, SKIMS, Spanx, Vuori)

$3–5T agentic commerce projected by 2030

AGP · UCP · AP2 · MCP · A2A

whoever owns ground truth owns the next economy.

APPENDIX · 02 — WHY NOT THE LABS

they've already *said* they won't.

> OPENAI ON INSTANT CHECKOUT
>
> *"product results are organic and unsponsored, ranked purely on relevance. merchant payments don't influence ChatGPT's product results."*

the moment they pick winners, they lose *neutrality* — and neutrality is the entire reason consumers trust agents to shop. they need an external source of ground truth. right now, they're scraping it from Wirecutter and Reddit — exactly the polluted data we started Giftly to replace.

we're not their competitor. we're the missing infrastructure.

APPENDIX · 03 — WHY NOT MERCOR

wrong customer. wrong economics. wrong data.

wrong customer	Mercor sells to AI labs (RLHF for capability training). brand-side affiliate teams are a different buyer with different sales cycle, price point, and metric.
wrong economics	Mercor pays experts hourly. we get eval data as a byproduct of brand-funded creator commissions — zero labor cost to us. capital advantage flips.
wrong data	paid evaluators with no audience produce judgment without behavioral validation. that's structurally what gameable reviews already are.

the real threat is a commerce platform with brand relationships. but the trust layer has to be neutral — same conflict the labs have.

APPENDIX · 04 — WHY US

we don't have an attention-economics background. we have an *AI capabilities* one.

01 the past year: building agent *harnesses* for Mandarins and ARC.

02 research toolkit: *Bayesian inverse planning*, human-AI interaction, decision-making under limited noisy human input.

03 ethan: *8 publications*, patent pending. builds in a day what takes most people a week.

04 most people in this space don't know this toolkit exists.

APPENDIX · 05 — THE TEAM

three founders. one decision rule.

ENGINEERING

Ethan Zhou

Berkeley · 8 publications · patent pending · ARC + Mandarins agent harness work.

"builds in a day what takes most people a week."

STRATEGY + BRAND

Samarjit Deshmukh

Dartmouth · competitive debate · creator outreach + brand conversations + strategy.

"the one in the room when a brand has hard questions."

DESIGN + CURATION

Armaan Priyadarshan

Dartmouth · math, art, UI/UX · 100+ nights with Samarjit in the CS basement, 11 projects, 3 hackathons.

"the most particular about details most people wouldn't notice."

product → Armaan. strategy → Samarjit. engineering → Ethan. anything cross-cutting, all three.

APPENDIX · 06 — HOW WE GOT HERE

we tried to fix our own *dandruff*.

Samarjit and I both have seborrheic dermatitis. one late night working on Mandarins, debating dandruff shampoos — we'd both spent *$60+* on Amazon-floated products that didn't work. neither of us could tell what was good.

we asked: how did these products even reach us? we built a side project — gifting between brands and creators with no strings, content only if it was real. *10–15 creator conversations* later, we launched.

APPENDIX · 07 — THE PATH

from now to series A in *15 months*.

pre-PearX *now → Jul 2026*	v1 product live: creator + brand portals, gifting workflow, MCP server v1 ingesting signal. eval study v0.	**2,500** contributors **200** brands · ~$3B valuation **500+** matches completed
PearX demo day *Oct 2026*	ChatGPT App + Perplexity Plugin live. 5+ agent surface integrations. eval study v1 ≥20% improvement on public rubric.	**12,000** contributors (creators + experts + power-users) **1,000** brands · ≥$10B **$600k/mo** GMV · **$150k** MRR
month 9 *Apr 2027*	Claude + Gemini integrations. eval paper accepted at workshop. seed round closed $5–8M.	**30,000** contributors **1,500** brands **$2M/mo** GMV · **$750k** MRR
month 15 *Oct 2027 — series A*	7+ agent surface integrations. featured in OpenAI / Anthropic / Perplexity case studies. enterprise contracts ($50k+/yr).	**80,000+** contributors **4,000+** brands **$3M+** ARR · **$50M+** annualized GMV

APPENDIX · 08 — WHAT WE NEED

$1.2M. 15 months. zero engineering hires.

BUCKETS

survival	31%	*$375k*
scaling	58%	*$700k*
legal / G&A	4%	*$50k*
reserve	6%	*$75k*

SCALING BLOCK

brand acquisition (7 → 1,500+)	*$275k*
agent integration + eval moat	*$200k*
contributor network (50 → 25,000)	*$130k*
platform + data infra	*$50k*
content + credibility (research paper, PR)	*$45k*

3 founders × $50k × 15 mo + 1 sr. brand BD lead (mo. 3) + 1 contributor ops lead (mo. 7). founders build. capital scales the network.

APPENDIX · 09 — PMF

early *signals*, not yet PMF.

STAGE 1

wedge PMF

brand × creator marketplace works — CPA contracts close, commissions flow, both sides return.

→ early signals

STAGE 2

data PMF

eval dataset large + structured enough to answer interesting product-quality queries.

→ ~1k+ structured records

STAGE 3

agent PMF

agents query API in production. turning it off would visibly degrade their recommendations.

→ outreach in motion

DEFENSIBLE EARLY SIGNALS

- *0.67%* brand cold-conversion (7/1,043) · solid for B2B
- *Apollo Neuro* inbound CPA offer — willingness-to-pay, unprompted
- *$300M* partner-tier resonance in 1 week
- *1,043* brands in active outreach

LEADING INDICATORS WE'LL PROVE OUT

- first revshare check (closes unit-economic loop)
- first inbound brand inquiry
- creator-to-creator referrals
- first agent partner technical conversation

in active outreach: Firecrawl · Browserless · OneOff · doji · PayOS · Skyfire.

Mercor for trust.

Mercor built the human-*expertise* layer that frontier AI labs run on for capability training. Giftly is building the human-*trust* layer that frontier AI agents run on for commerce decisions.

same shape — humans as infrastructure for AI. different customer. different signal. market an order of magnitude larger.

---

by year 5: every brand selling through agent surfaces, every human with verifiable expertise, every agent shopping decision — runs through Giftly.